

DP-900

Real Exam Questions & Answers

Latest Exam Questions - 2023
All Topics Covered
Added New Questions



Get Certified Today!

New Course Covered

Subscribe



Question 201

CertyIQ

Your company is designing a data store for internet-connected temperature sensors.

The collected data will be used to analyze temperature trends.

Which type of data store should you use?

- A. relational
- B. time series
- C. graph
- D. columnar

Explanation:

Correct Answer: B

Time series data is a set of values organized by time. Time series databases typically collect large amounts of data in real time from a large number of sources.

Updates are rare, and deletes are often done as bulk operations. Although the records written to a time-series database are generally small, there are often a large number of records, and total data size can grow rapidly.

Link for time series

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/analytical-data-stores>

Telemetry and time series databases are an append-only collection of objects. Telemetry databases efficiently index data in a variety of column stores and in-memory structures, making them the optimal choice for storing and analyzing vast quantities of telemetry and time series data.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

Question 202

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Copying data to Azure Data Lake Storage from another Azure region results in extra cost.	<input type="radio"/>	<input type="radio"/>
You can use blob, table, and file storage in the same Azure Storage account.	<input type="radio"/>	<input type="radio"/>
You implement Azure Data Lake Storage by creating an Azure Storage account.	<input type="radio"/>	<input type="radio"/>

Answer Area

	Statements	Yes	No
Suggested Answer:	Copying data to Azure Data Lake Storage from another Azure region results in extra cost.	<input checked="" type="radio"/>	<input type="radio"/>
	You can use blob, table, and file storage in the same Azure Storage account.	<input checked="" type="radio"/>	<input type="radio"/>
	You implement Azure Data Lake Storage by creating an Azure Storage account.	<input checked="" type="radio"/>	<input type="radio"/>

Explanation:

Correct Answer: YYY

Reference:

<https://docs.microsoft.com/en-us/azure/data-lake-store/data-lake-store-get-started-portal>

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-overview>

<https://azure.microsoft.com/en-us/pricing/details/bandwidth/>

Question 203

CertyIQ

HOTSPOT -

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

	▼
graph.	
table.	
partition key.	
document.	

Answer Area

Suggested Answer:

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

	▼
graph.	
table.	
partition key.	
document.	

Explanation:

Correct Answer:

Graph = Gremlin

Reference : <https://azure.microsoft.com/en-us/blog/a-technical-overview-of-azure-cosmos-db/>

API - Gremlin // Container is projected as -Graph // Item is projected as -Nodes and Edges

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/create-graph-gremlin-console>

Question 204

CertyIQ

At which two levels can you set the throughput for an Azure Cosmos DB account? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

A. database

B. item

C. container

D. partition

Explanation:

Correct Answer: AC

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/set-throughput>

Question 205

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	<input type="radio"/>	<input type="radio"/>
Azure Table storage supports multiple write regions.	<input type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple read replicas.	<input type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple write regions.	<input type="radio"/>	<input type="radio"/>

Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	<input checked="" type="radio"/>	<input type="radio"/>
Azure Table storage supports multiple write regions.	<input type="radio"/>	<input checked="" type="radio"/>
The Azure Cosmos DB Table API supports multiple read replicas.	<input checked="" type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple write regions.	<input checked="" type="radio"/>	<input type="radio"/>

Suggested Answer:

Explanation:

Correct Answer:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/high-availability>

Question 206

CertyIQ

What is a characteristic of non-relational data?

A. no indexes

B. self-describing entities

- C. a well-defined schema
- D. no unique key values

Explanation:

Correct Answer: B

<https://www.coursera.org/lecture/explore-core-data-concepts-microsoft-azure/explore-the-characteristics-of-non-relational-data-nP41V>

Question 207

CertyIQ

HOTSPOT -

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

	▼
A hierarchical namespace	
Azure Synapse Link	
PolyBase	
The Data Movement Service (DMS)	

is an Azure Synapse Analytics feature that retrieves data from Azure Blob storage by using Transact-SQL queries.

Answer Area

Suggested Answer:

	▼
A hierarchical namespace	
Azure Synapse Link	
PolyBase	
The Data Movement Service (DMS)	

is an Azure Synapse Analytics feature that retrieves data from Azure Blob storage by using Transact-SQL queries.

Explanation:

Correct Answer: Polybase

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/polybase/polybase-guide?view=sql-server-ver15>

Question 208

CertyIQ

You need to gather real-time telemetry data from a mobile application.

Which type of workload describes this scenario?

- A. Online Transaction Processing (OLTP)
- B. batch

C. massively parallel processing (MPP)

D. streaming

Explanation:

Correct Answer: D

as real-time is always Streaming.. OLTP is delayed upto 30 sec

Reference:

<https://docs.microsoft.com/en-in/azure/azure-monitor/overview>

Question 209

CertyIQ

You have a dedicated SQL pool in Azure Synapse Analytics that is only used actively every night for eight hours.

You need to minimize the cost of the dedicated SQL pool as much as possible during idle times. The solution must ensure that the data remains intact.

What should you do on the dedicated SQL pool?

A. Scale down the data warehouse units (DWUs).

B. Pause the pool.

C. Create a user-defined restore point.

D. Delete the pool

Explanation:

Correct Answer: B Pause the pool

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-manage-compute-overview>

Question 210

CertyIQ

Which Azure Data Factory component initiates the execution of a pipeline?

A. a control flow

B. a trigger

C. a parameter

D. an activity

Explanation:

Correct Answer: B

Trigger.

Pipeline runs are typically instantiated by passing arguments to parameters that you define in the pipeline. You can execute a pipeline either manually or by using a trigger.

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipeline-execution-triggers#trigger-execution>

Question 211

CertyIQ

Your company has a reporting solution that has paginated reports. The reports query a dimensional model in a data warehouse.

Which type of processing does the reporting solution use?

- A. stream processing
- B. batch processing
- C. Online Analytical Processing (OLAP)
- D. Online Transaction Processing (OLTP)

Explanation:

Correct Answer: C

DWH --> OLAP

data warehouse is for OLAP, DBMS is for OLTP

Reference:

<https://datawarehouseinfo.com/how-does-oltp-differ-from-olap-database/>

Question 212

CertyIQ

DRAG DROP -

Match the types of activities to the appropriate Azure Data Factory activities.

To answer, drag the appropriate activity type from the column on the left to its Data Factory activity on the right. Each activity type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Activity Types	Answer Area	
Control		Copy
Data movement		Mapping data flow
Data transformation		Lookup

	Activity Types	Answer Area	
Suggested Answer:	Control	Data movement	Copy
	Data movement	Data transformation	Mapping data flow
	Data transformation	Control	Lookup

Explanation:

Correct Answer:

Box 1: Data movement -

Box 2: Data transformation -

A pipeline could contain a set of activities that ingest and clean log data, and then kick off a mapping data flow to analyze the log data.

Box 3: Control -

Lookup Activity is a control flow activity.

Lookup Activity can be used to read or look up a record/ table name/ value from any external source. This output can further be referenced by succeeding activities.

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipelines-activities#data-transformation-activities>

Question 213

CertyIQ

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

A. denormalized data

B. heavy writes and moderate reads

C. light writes and heavy reads

D. schema on write

E. schema on read

F. normalized data

Explanation:

Correct Answer: BDF

B: Transactional data tends to be heavy writes, moderate reads.

D: Typical traits of transactional data include: schema on write, strongly enforced

F: Transactional data tends to be highly normalized.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>

Question 214

CertyIQ

Which two activities can be performed entirely by using the Microsoft Power BI service without relying on Power BI Desktop? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

A. report and dashboard creation

B. report sharing and distribution

C. data modeling

D. data acquisition and preparation

Explanation:

Correct Answer: AB

A. reports/dashboard

B. sharing

according to this link: <https://docs.microsoft.com/en-us/power-bi/fundamentals/service-service-vs-desktop>

Question 215

CertyIQ

HOTSPOT -

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

In Azure Data Factory, you can use

a control flow

a dataset

a linked service

an integration runtime

activities that depend on the output of other pipeline activities.

Answer Area

Suggested Answer:

In Azure Data Factory, you can use

	▼
a control flow	
a dataset	
a linked service	
an integration runtime	

 to orchestrate pipeline

activities that depend on the output of other pipeline activities.

Explanation:

Correct Answer: Control flow (<https://docs.microsoft.com/en-us/azure/data-factory/introduction>)

Control flow is an orchestration of pipeline activities that includes chaining activities in a sequence, branching, defining parameters at the pipeline level, and passing arguments while invoking the pipeline on-demand or from a trigger. It also includes custom-state passing and looping containers, that is, For-each iterators.

Question 216

CertyIQ

You need to develop a solution to provide data to executives. The solution must provide an interactive graphical interface, depict various key performance indicators, and support data exploration by using drill down.

What should you use in Microsoft Power BI?

- A. a view
- B. a report**
- C. a dataflow
- D. Microsoft Power Apps

Explanation:

Correct Answer: B

Reference:

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards>

<https://docs.microsoft.com/en-us/power-bi/visuals/power-bi-visualization-kpi>

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-drill>

Question 217

CertyIQ

Which two Azure services can be used to provision Apache Spark clusters? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Azure Time Series Insights
- B. Azure HDInsight**

C. Azure Databricks

D. Azure Log Analytics

Explanation:

Correct Answer: BC

<https://docs.microsoft.com/en-us/azure/hdinsight/spark/apache-spark-overview>

Question 218

CertyIQ

You have a quality assurance application that reads data from a data warehouse.

Which type of processing does the application use?

A. Online Transaction Processing (OLTP)

B. batch processing

C. Online Analytical Processing (OLAP)

D. stream processing

Explanation:

Correct Answer: C

Quality assurance application focuses on improving the software development process and making it efficient and effective as per quality standards.

OLAP applies complex queries to large amounts of historical data, aggregated from OLTP databases and other sources, for data mining, analytics, and business intelligence projects.

Moreover, the purpose of OLAP ----> Plan, solve problems, support decisions, discover hidden insights.

According to the question, the application is reading data from the data warehouse for quality assurance which is clearly - discovering insights from the data. Thus, the answer is OLAP.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>

Question 219

CertyIQ

Which three objects can be added to a Microsoft Power BI dashboard? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

A. a report page

B. a Microsoft PowerPoint slide

C. a visualization from a report

D. a dataflow

E. a text box

Explanation:

Correct Answer: ACE

A,C,E

<https://docs.microsoft.com/en-us/power-bi/create-reports/service-dashboard-pin-live-tile-from-report>

<https://docs.microsoft.com/en-us/power-bi/create-reports/service-dashboard-add-widget>

Question 220

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input type="radio"/>	<input type="radio"/>

Answer Area

Suggested Answer:

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	<input checked="" type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input checked="" type="radio"/>	<input type="radio"/>

Explanation:

Correct Answer: Yes , No , yes

Dashboard is associate for a single workspace .

Dashboard visualization from different datasets and reports.

Dashboard visualization from other tools example: Excel

"ONE dashboard..."

- is associated with a single workspace
- can display visualizations from many different datasets
- can display visualizations from many different reports
- can display visualizations pinned from other tools (for example, Excel)"

Reference: <https://docs.microsoft.com/en-us/power-bi/fundamentals/service-basic-concepts>

Question 221

CertyIQ

Which Azure Data Factory component provides the compute environment for activities?

- A. SSIS packages
- B. an integration runtime**
- C. a control flow
- D. a pipeline

Explanation:

Correct Answer: B

The Integration Runtime (IR) is the compute infrastructure used by Azure Data Factory to provide the following data integration capabilities across different network environments:

- ⇒ Data Flow
- ⇒ Data movement
- ⇒ Activity dispatch
- ⇒ SSIS package execution

Reference: <https://docs.microsoft.com/en-us/azure/data-factory/concepts-integration-runtime>

Question 222

CertyIQ

HOTSPOT -

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

A Microsoft Power BI

	▼
dashboard	
interactive report	
paginated report	
subscription	

enables users to create highly formatted,

fixed-layout documents optimized for printing and archiving.

Answer Area

Suggested Answer:

A Microsoft Power BI

▼
dashboard
interactive report
paginated report
subscription

enables users to create highly formatted,

fixed-layout documents optimized for printing and archiving.

Explanation:

Correct Answer:

Paginated Reports in Power BI now allows users to generate these fixed-layout documents optimized for printing and archiving, such as PDF and Word files.

These document-style reports with visualizations that provide additional control, like which tables expand horizontally and vertically to display all their data and continue from page to page as needed.

Reference:

<https://powerbi.microsoft.com/en-us/blog/announcing-paginated-reports-in-power-bi-general-availability/>

Question 223

CertyIQ

What are two uses of data visualization? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Represent trends and patterns over time
- B. Implement machine learning to predict future values
- C. Communicate the significance of data
- D. Enforce business logic across reports

Explanation:

Correct Answer: AC

Data visualization is a key component in being able to gain insight into your data. It helps make big and small data easier for humans to understand. It also makes it easier to detect patterns, trends, and outliers in groups of data.

Data visualization brings data to help you find key business insights quickly and effectively.

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/spark/apache-spark-data-visualization>

Question 224

CertyIQ

You need to use Transact-SQL to query files in Azure Data Lake Storage Gen 2 from an Azure Synapse Analytics data warehouse.

What should you use to query the files?

- A. Azure Functions
- B. Microsoft SQL Server Integration Services (SSIS)
- C. PolyBase
- D. Azure Data Factory

Explanation:

Correct Answer: C

Polybase is the correct answerd,

<https://docs.microsoft.com/en-gb/learn/modules/explore-data-storage-processing-azure/3-explore-azure-synapse-analytics>

"Use SQL pools in Synapse Analytics for the following scenarios:

Complex reporting. You can use the full power of Transact-SQL to run complex SQL statements that summarize and aggregate data.

Data ingestion. PolyBase enables you to retrieve data from many external sources and convert it into a tabular format. You can reformat this data and save it as tables and materialized views in Azure Synapse."

Question 225

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
You can copy a dashboard between Microsoft Power BI workspaces.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input type="radio"/>	<input type="radio"/>

Suggested Answer:

Answer Area

Statements	Yes	No
You can copy a dashboard between Microsoft Power BI workspaces.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input checked="" type="radio"/>	<input type="radio"/>

Explanation:

Correct Answer:

Box 1: No -

You can duplicate a dashboard. The duplicate ends up in the same Power BI workspace.

There is no current functionality that allows you to move reports from one workspace to another.

Box 2: No -

Box 3: Yes -

Reference:

<https://docs.microsoft.com/en-us/power-bi/connect-data/service-datasets-across-workspaces>

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards>

<https://powerbi.microsoft.com/en-us/excel-and-power-bi/>

Question 226

CertyIQ

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

A. denormalized data

B. heavy writes and moderate reads

C. light writes and heavy reads

D. schema defined in a database

E. schema defined when reading unstructured data from a database

F. normalized data

Explanation:

Correct Answer: BDF

B: Transactional data tends to be heavy writes, moderate reads.

D: Typical traits of transactional data include: schema on write, strongly enforced. The schema is defined in a database.

F: Transactional data tends to be highly normalized.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>

Question 227

CertyIQ

What is the primary purpose of a data warehouse?

A. to provide answers to complex queries that rely on data from multiple sources

B. to provide transformation services between source and target data stores

C. to provide read-only storage of relational and non-relational historical data

D. to provide storage for transactional line-of-business (LOB) applications

Explanation:

Correct Answer: A

this is the main purpose of a data warehouse.

With C you could just create read-only replica's.

In the textbook DP-900T00-A-Microsoft Azure Data Fundamentals Module 4 Describe modern data warehousing:

"A data warehouse gathers data from many different sources within an organization...The focus of a data warehouse is to provide answers to complex queries..."

Question 228

CertyIQ

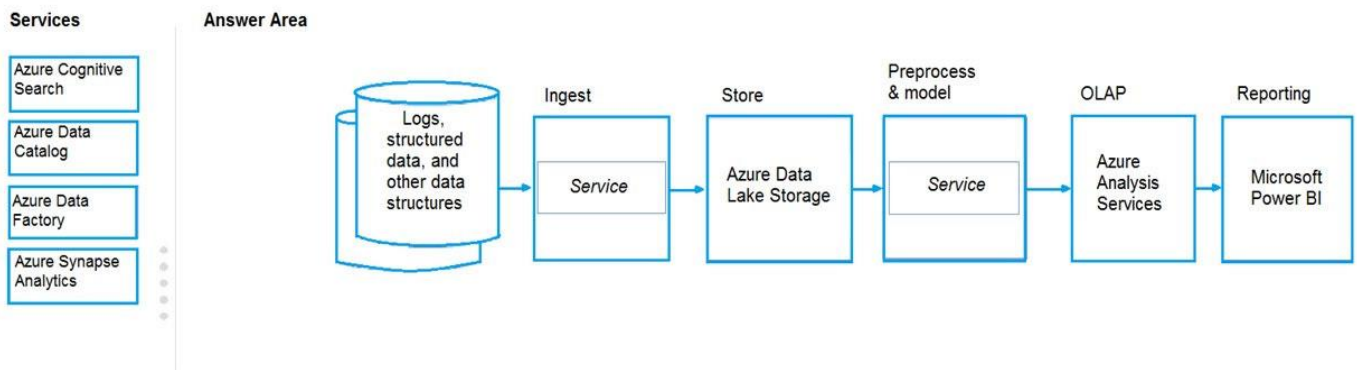
DRAG DROP -

Match the Azure services to the appropriate locations in the architecture.

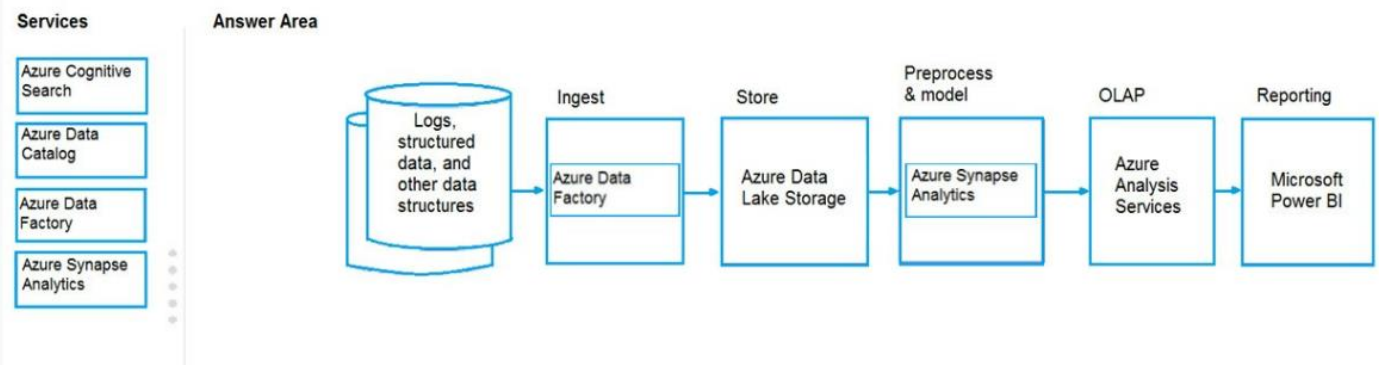
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:



Suggested Answer:



Explanation:

Correct Answer:

Box Ingest: Azure Data Factory -

You can build a data ingestion pipeline with Azure Data Factory (ADF).

Box Preprocess & model: Azure Synapse Analytics

Use Azure Synapse Analytics to preprocess data and deploy machine learning models.

Reference:

<https://docs.microsoft.com/en-us/azure/machine-learning/how-to-data-ingest-adf>

<https://docs.microsoft.com/en-us/azure/machine-learning/team-data-science-process/sqldw-walkthrough>

Question 229

CertyIQ

DRAG DROP -

Match the types of workloads to the appropriate scenarios.

To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Workload Types

Answer Area

Batch

Workload type

Analyzing historical data containing web traffic collected during the previous year.

Streaming

Workload type

Classifying images that were uploaded last month.

Workload type

Tracking in real time how many people are currently using a website.

	Workload Types	Answer Area
Suggested Answer:	Batch	Batch Analyzing historical data containing web traffic collected during the previous year.
	Streaming	Batch Classifying images that were uploaded last month.
		Streaming Tracking in real time how many people are currently using a website.

Explanation:

Correct Answer:

Box 1: Batch -

The batch processing model requires a set of data that is collected over time while the stream processing model requires data to be fed into an analytics tool, often in micro-batches, and in real-time.

The batch Processing model handles a large batch of data while the Stream processing model handles individual records or micro-batches of few records.

In Batch Processing, it processes over all or most of the data but in Stream Processing, it processes over data on a rolling window or most recent record.

Box 2: Batch -

Question 230

CertyIQ

HOTSPOT -

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

In a data warehousing workload, data

from a single source is distributed to multiple locations
from multiple sources is combined in a single location
is added to a queue for multiple systems to process
is used to train machine learning models

Suggested Answer:

Answer Area

In a data warehousing workload, data

from a single source is distributed to multiple locations
from multiple sources is combined in a single location
is added to a queue for multiple systems to process
is used to train machine learning models

Explanation:

Correct Answer:

Note: The data warehouse workload encompasses:

- ⇒ The entire process of loading data into the warehouse
- ⇒ Performing data warehouse analysis and reporting
- ⇒ Managing data in the data warehouse
- ⇒ Exporting data from the data warehouse

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-management>

Question 231

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory	<input type="radio"/>	<input type="radio"/>
Azure Data Factory pipelines can execute other pipelines	<input type="radio"/>	<input type="radio"/>
A processing step within an Azure Data Factory pipeline is an activity	<input type="radio"/>	<input type="radio"/>

Answer Area

Suggested Answer:

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory	<input type="radio"/>	<input checked="" type="radio"/>
Azure Data Factory pipelines can execute other pipelines	<input checked="" type="radio"/>	<input type="radio"/>
A processing step within an Azure Data Factory pipeline is an activity	<input checked="" type="radio"/>	<input type="radio"/>

Explanation:

Correct Answer:

Box 1: No -

A pipeline is a logical grouping of activities that together perform a task.

Box 2: Yes -

You can construct pipeline hierarchies with data factory.

Box 3: Yes -

A pipeline is a logical grouping of activities that together perform a task.

Reference:

<https://mrpaulandrew.com/2019/09/25/azure-data-factory-pipeline-hierarchies-generation-control/>

Question 232

CertyIQ

DRAG DROP -

Match the Azure services to the appropriate requirements.

To answer, drag the appropriate service from the column on the left to its requirement on the right. Each service may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Services

Azure Data Lake Storage

Azure SQL Database

Azure Synapse Analytics

Answer Area

Service

Output data to Parquet format

Service

Store data that is in Parquet format

Service

Persist a tabular representation of data that is stored in Parquet format

Suggested Answer:

Services

Azure Data Lake Storage

Azure SQL Database

Azure Synapse Analytics

Answer Area

Azure Synapse Analytics

Output data to Parquet format

Azure Data Lake Storage

Store data that is in Parquet format

Azure SQL Database

Persist a tabular representation of data that is stored in Parquet format

Explanation:

Correct Answer:

- Azure Synapse Analytics: Output data to parquet format
- Azure Data Lake Storage: Store data that is in parquet format
- Azure SQL Database: Persist a tabular representation of data that is stored in parquet format

Question 233

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure Synapse Analytics scales storage and compute independently	<input type="radio"/>	<input type="radio"/>
Azure Synapse Analytics can be paused to reduce compute costs	<input type="radio"/>	<input type="radio"/>
An Azure Synapse Analytics data warehouse has a fixed storage capacity	<input type="radio"/>	<input type="radio"/>

Suggested Answer:

Answer Area

Statements	Yes	No
Azure Synapse Analytics scales storage and compute independently	<input checked="" type="radio"/>	<input type="radio"/>
Azure Synapse Analytics can be paused to reduce compute costs	<input checked="" type="radio"/>	<input type="radio"/>
An Azure Synapse Analytics data warehouse has a fixed storage capacity	<input type="radio"/>	<input checked="" type="radio"/>

Explanation:

Correct Answer:

Box 1: Yes -

Compute is separate from storage, which enables you to scale compute independently of the data in your system.

Box 2: Yes -

You can use the Azure portal to pause and resume the dedicated SQL pool compute resources.

Pausing the data warehouse pauses compute. If your data warehouse was paused for the entire hour, you will not be charged compute during that hour.

Box 3: No -

Storage is sold in 1 TB allocations. If you grow beyond 1 TB of storage, your storage account will automatically grow to 2 TBs.

Reference:

<https://azure.microsoft.com/en-us/pricing/details/synapse-analytics/>

Question 234

CertyIQ

What should you use to build a Microsoft Power BI paginated report?

- A. Charticator
- B. Power BI Desktop
- C. the Power BI service
- D. Power BI Report Builder**

Explanation:

Correct Answer: D

Power BI Report Builder is the standalone tool for authoring paginated reports for the Power BI service.

Reference:

<https://docs.microsoft.com/en-us/power-bi/paginated-reports/paginated-reports-report-builder-power-bi>

Question 235

DRAG DROP -

Match the Azure services to the appropriate locations in the architecture.

To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.

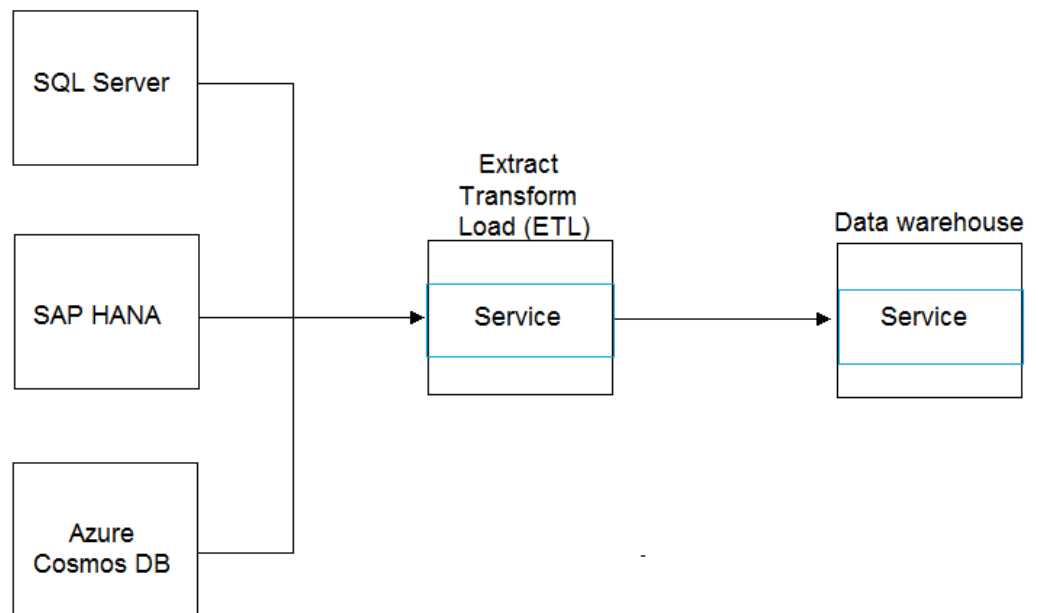
NOTE: Each correct match is worth one point.

Select and Place:

Services

- Azure Analysis Services
- Azure Data Factory
- Azure Table storage
- Azure Cosmos DB
- Azure Synapse Analytics

Answer Area

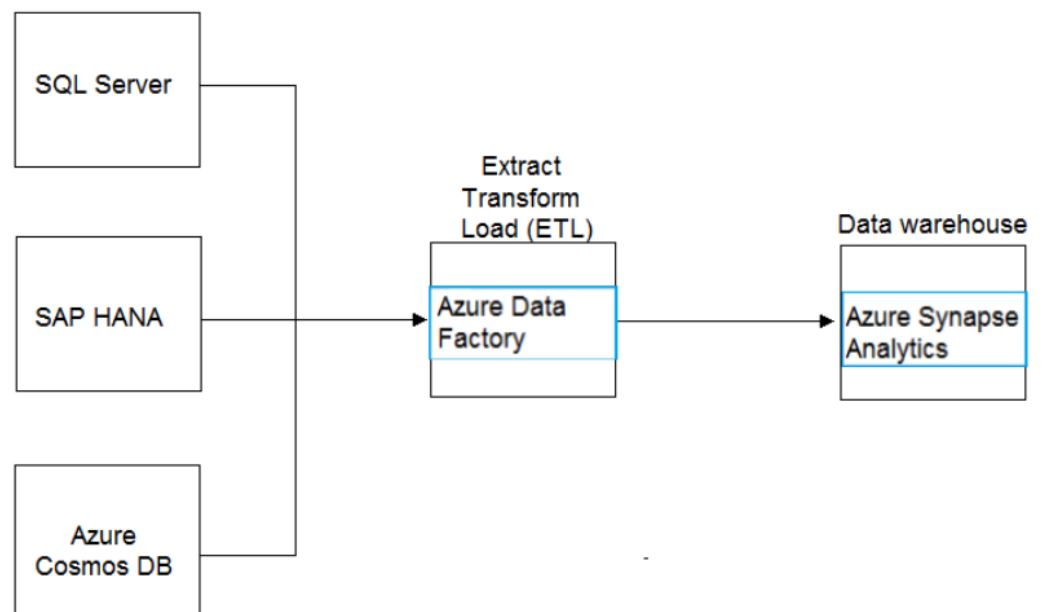


Suggested Answer:

Services

- Azure Analysis Services
- Azure Data Factory
- Azure Table storage
- Azure Cosmos DB
- Azure Synapse Analytics

Answer Area



Explanation:

Correct Answer:

Box 1: Azure Data factory -

Relevant Azure service for the three ETL phases are Azure Data Factory and SQL Server Integration Services (SSIS).

Box 2: Azure Synapse Analytics -

You can copy and transform data in Azure Synapse Analytics by using Azure Data Factory

Note: Azure Synapse Analytics connector is supported for the following activities:

⇒ Copy activity with supported source/sink matrix table

⇒ Mapping data flow

⇒ Lookup activity

⇒ GetMetadata activity

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/etl>

<https://docs.microsoft.com/en-us/azure/data-factory/connector-azure-sql-data-warehouse>

Question 236

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure Databricks is an Apache Spark-based analytics platform.	<input type="radio"/>	<input type="radio"/>
Azure Analysis Services is used for transactional workloads.	<input type="radio"/>	<input type="radio"/>
Azure Data Factory orchestrates data integration workflows.	<input type="radio"/>	<input type="radio"/>

Answer Area			
Suggested Answer:	Statements	Yes	No
	Azure Databricks is an Apache Spark-based analytics platform.	<input checked="" type="radio"/>	<input type="radio"/>
	Azure Analysis Services is used for transactional workloads.	<input type="radio"/>	<input checked="" type="radio"/>
	Azure Data Factory orchestrates data integration workflows.	<input checked="" type="radio"/>	<input type="radio"/>

Explanation:

Correct Answer:

Reference:

Question 237

CertyIQ

Which scenario is an example of a streaming workload?

- A. sending transactions that are older than a month to an archive
- B. sending transactions daily from point of sale (POS) devices
- C. sending telemetry data from edge devices
- D. sending cloud infrastructure metadata every 30 minutes

Explanation:

Correct Answer: C

Because telemetry data from edge devices need to be processed in real-time, not in batches.

Question 238

CertyIQ

HOTSPOT -

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Batch workloads

	process data in memory, row-by-row.
	collect and process data at most once a day.
	process data as new data is received in near real-time.
	collect data and then process the data when a condition is met.

Answer Area

Suggested Answer:

Batch workloads

	process data in memory, row-by-row.
	collect and process data at most once a day.
	process data as new data is received in near real-time.
	collect data and then process the data when a condition is met.

Explanation:

Correct Answer:

A) doesn't make any sense as there is no data processing in the background. Batch means to process a big volume of data in a particular time of the day (e.g. in the night with idle VMs)

B) You can once per day if you define it as such but in any case you are limited by "one day". So it's no.

C) Nonsense. Batch data is not a design for RT/streaming data OTA. You would use a different solution as data bricks. Reason that data bricks service exists.

D) Correct. You can define your condition. Once it's met, you can execute it.

Question 239

CertyIQ

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	<input type="radio"/>	<input type="radio"/>
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	<input type="radio"/>	<input type="radio"/>
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	<input type="radio"/>	<input type="radio"/>

Answer Area

	Statements	Yes	No
Suggested Answer:	Processing salary payments once a month is an example of a batch workload.	<input checked="" type="radio"/>	<input type="radio"/>
	A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	<input checked="" type="radio"/>	<input type="radio"/>
	A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	<input type="radio"/>	<input checked="" type="radio"/>

Explanation:

Correct Answer:

Option C says it sends the "readingS" once a day, which makes me think it has multiple readings and wait to send them all once a day. So with given information I'm more inclined to Yes, Yes, No.

Question 240

CertyIQ

A bar chart showing year-to-date sales by region is an example of which type of analytics?

- A. predictive
- B. prescriptive
- C. descriptive**
- D. diagnostic

Explanation:

Correct Answer: C

Nothing is being diagnosed here, but something is being described. So I would say C

End of Part 6



Want all PDFs of DP900?

Please drop us mail at: - certyiq@gmail.com

Please find the videos of this **AZ-900/AI-900/AZ-305/AZ-104 /DP-900/ SC-900** and other Microsoft exam series on

CertyIQ Official YouTube channel (**FREE PDFs**): -

Please **Subscribe** to CertyIQ YouTube Channel to get notified for latest exam dumps by clicking on the below image, it will redirect to the **CertyIQ** YouTube page.



Connect with us @ [LinkedIn](#) [Telegram](#)

Contact us for other dumps: - certyiqpdf@gmail.com

For any other enquiry, please drop us a mail at

certyiqofficial@gmail.com